

1.	Prerequisite .....	2
2.	Installing Confluent Kafka (Local) – 60 Minutes .....	3
3.	Workflow using KSQL - CLI – 90 Minutes.....	19

### 1. Prerequisite

#### Scenario 1 : Using VM

Refer any tutorial in the web to configure Centos VM using VM Player or Workstation in your laptop.

Start the VM using VM player and Logon to the server using telnet or directly in the VM console. Enter the root credentials to logon.

#### Scenario 2: Using Docker

All the necessary software should be in the /Software folder. If its not there, ensure to copy it using winscp.exe from the windows desktop to /Software folder. You can create /Software folder using mkdir /Software.

The following instruction will create a network and bind to the container, ckafkao. Replace the -v parameter with any of the folder in your Host machine.

```
#docker network create --driver bridge spark-net
```

```
#docker run --name ckafkao --hostname ckafkao -p 9094:9092 -p 8086:8081 -p  
2184:2181 -p 9031:9021 -p 8098:8088 -i -t --privileged --network spark-net -v  
/Users/henrypotsangbam/Documents/Docker:/opt centos:7 /usr/sbin/init
```

## 2. Installing Confluent Kafka (Local) – 60 Minutes

Demonstrates both the basic and most powerful capabilities of Confluent Platform, including using Control Center for topic management and event stream processing using KSQL. In this quick start you create Apache Kafka® topics, use Kafka Connect to generate mock data to those topics, and create KSQL streaming queries on those topics. You then go to Control Center to monitor and analyze the streaming queries.

You need to install java before installing zookeeper and Kafka.

Installing Java

```
#tar -xvf jdk-8u45-linux-x64.tar.gz -C /opt
```

Set in the path variable and JAVA\_HOME

[ex:

```
export JAVA_HOME=/opt/jdk
```

```
export PATH=$PATH:$JAVA_HOME/bin
```

]

Include in the profile as follow

```
[root@tos opt]# more ~/.bashrc
# .bashrc

# User specific aliases and functions

alias rm='rm -i'
alias cp='cp -i'
alias mv='mv -i'

export JAVA_HOME=/opt/jdk1.8.0_45

export PATH=$PATH:$JAVA_HOME/bin
# Source global definitions
if [ -f /etc/bashrc ]; then
    . /etc/bashrc
fi
[root@tos opt]#
```

### Installing a Kafka Broker

The following example installs Confluence Kafka in /apps.

### Installing and Configuring Confluent CLI

Inflate the confluent kafka compress file as shown below:

```
#tar -xvf confluent-5.5.1-2.12.tar -C /apps
```

Rename the folder.

```
#mv /apps/confluent* /apps/confluent
```

Set the environment variable for the Confluent Platform directory (<path-to-confluent>).

## 5 Kafka – Dev Ops

```
export CONFLUENT_HOME=/apps/confluent
```

```
(base) [root@tos confluent]# pwd
/apps/confluent
(base) [root@tos confluent]# ls
bin confluent etc legal lib logs README share src
(base) [root@tos confluent]#
```

Set your PATH variable:

```
# vi ~/.bashrc
```

```
export PATH=/apps/confluent/bin:${PATH};
```

```
if [ -f "/apps/anaconda3/etc/profile.d/conda.sh" ]; then
    . "/apps/anaconda3/etc/profile.d/conda.sh"
else
    export PATH="/apps/anaconda3/bin:${PATH}"
fi
fi
unset __conda_setup
# <<< conda initialize <<<

export JAVA_HOME=/apps/jdk
export PATH=$JAVA_HOME/bin:${PATH}:${SCALA_HOME/bin}
export PATH=/apps/confluent/bin:${PATH};
-- INSERT --
```

After decompressing the file. You should have the following directories:

```
(base) [root@tos confluent]#  
(base) [root@tos confluent]# pwd  
/apps/confluent  
(base) [root@tos confluent]# ls -ltr  
total 16  
drwxr-xr-x.  3 life life   21 Jun  5 10:11 lib  
drwxr-xr-x.  7 life life  106 Jun  5 10:42 share  
drwxr-xr-x. 23 life life 4096 Jun  5 10:42 etc  
drwxr-xr-x.  3 life life 4096 Jun  5 10:42 bin  
drwxr-xr-x.  2 life life  178 Jun  5 11:17 src  
-rw-r--r--.  1 life life  871 Jun  5 11:17 README  
drwxr-xr-x.  2 root root 4096 Jul  7 02:01 logs  
(base) [root@tos confluent]#
```

Install the [Kafka Connect Datagen](#) source connector using the Confluent Hub client. This connector generates mock data for demonstration purposes and is not suitable for production. [Confluent Hub](#) is an online library of pre-packaged and ready-to-install extensions or add-ons for Confluent Platform and Kafka.

```
#confluent-hub install --no-prompt confluentinc/kafka-connect-datagen:latest
```

```
(base) [root@tos ~]# cd /apps
(base) [root@tos apps]# confluent-hub install --no-prompt confluentinc/kafka-connect-datagen:latest
Running in a "--no-prompt" mode
Implicit acceptance of the license below:
Apache License 2.0
https://www.apache.org/licenses/LICENSE-2.0
Downloading component Kafka Connect Datagen 0.1.3, provided by Confluent, Inc. from Confluent Hub and installing into /apps/confluent/share/confluent-hub-components
Adding installation directory to plugin path in the following files:
  /apps/confluent/etc/kafka/connect-distributed.properties
  /apps/confluent/etc/kafka/connect-standalone.properties
  /apps/confluent/etc/schema-registry/connect-avro-distributed.properties
  /apps/confluent/etc/schema-registry/connect-avro-standalone.properties
  /tmp/confluent.8A2Ii7O4/connect/connect.properties

Completed
(base) [root@tos apps]#
```

Start Confluent Platform using the Confluent CLI `confluent local start` command. This command starts all of the Confluent Platform components; including Kafka, ZooKeeper, Schema Registry, HTTP REST Proxy for Kafka, Kafka Connect, KSQL, and Control Center.

```
#export CONFLUENT_CURRENT=/opt/data/ckafka
```

```
#confluent local services start
```

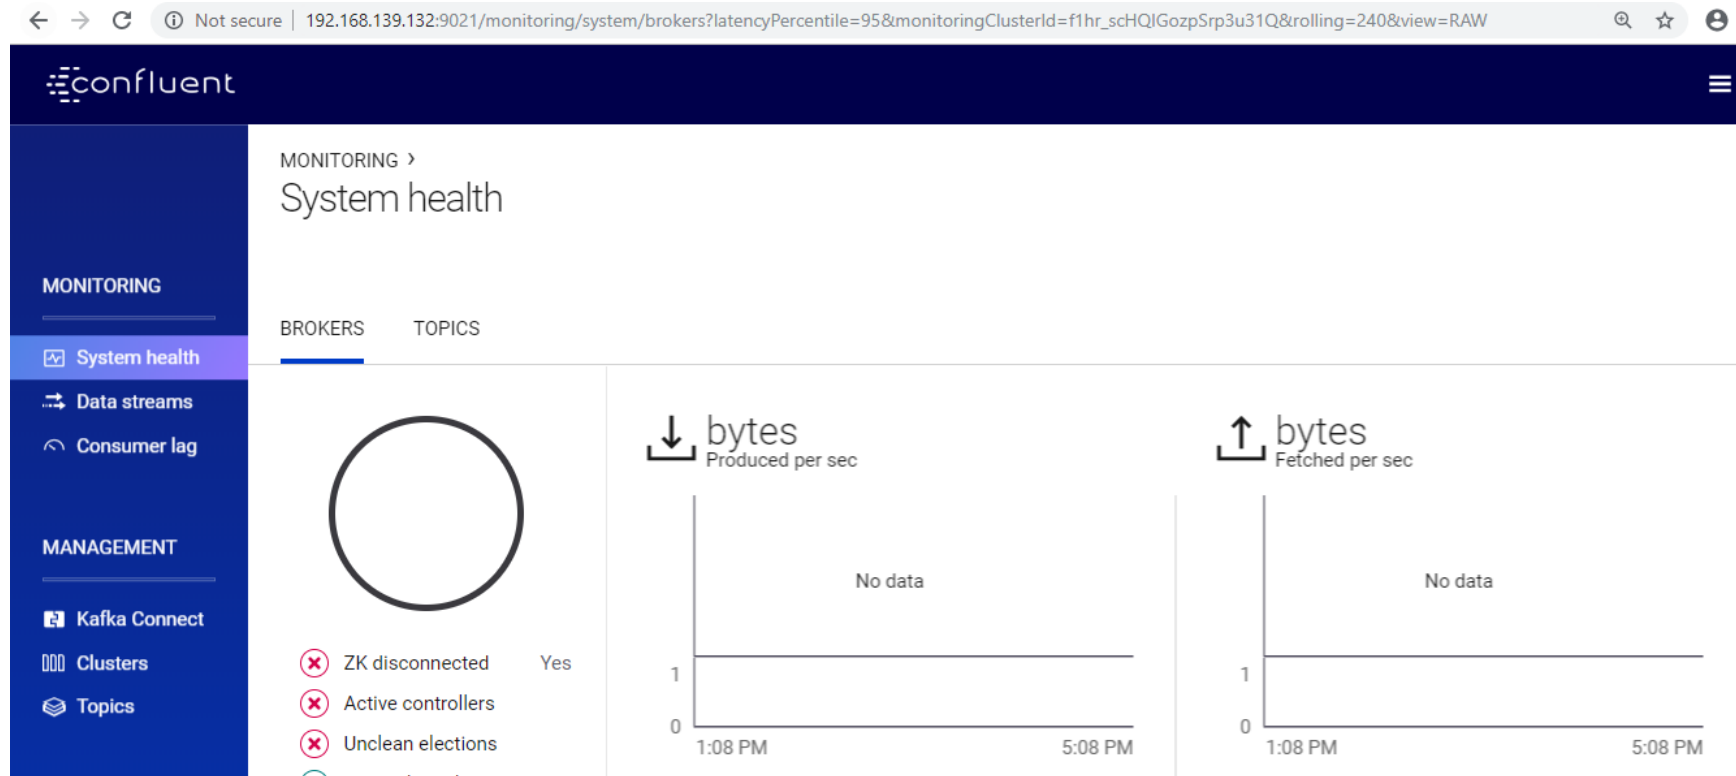
```
(base) [root@tos bin]# confluent start
This CLI is intended for development only, not for production
https://docs.confluent.io/current/cli/index.html

Using CONFLUENT_CURRENT: /tmp/confluent.8A2Ii704
Starting zookeeper
zookeeper is [UP]
Starting kafka
kafka is [UP]
Starting schema-registry
schema-registry is [UP]
Starting kafka-rest
kafka-rest is [UP]
Starting connect
connect is [UP]
Starting ksql-server
ksql-server is [UP]
Starting control-center
control-center is [UP]
(base) [root@tos bin]#
```

Navigate to the Control Center web interface at <http://localhost:9021/>.



## 9 Kafka – Dev Ops



## Install a Kafka Connector and Generate Sample Data

In this step, you use Kafka Connect to run a demo source connector called kafka-connect-datagen that creates sample data for the Kafka topics pageviews and users.

Run one instance of the [Kafka Connect Datagen](#) connector to produce Kafka data to the `pageviews` topic in AVRO format.

Management → **Add connector**. Or Connectors → Add Connector

Find the DatagenConnector tile and click **Connect**.

Name the connector `datagen-pageviews`. After naming the connector, new fields appear. Scroll down and specify the following configuration values:

- Tasks max : 1
- In the **Key converter class** field, type `org.apache.kafka.connect.storage.StringConverter`.
- In the **kafka.topic** field, type `pageviews`.
- In the **max.interval** field, type `100`.
- In the **iterations** field, type `1000000000`.
- In the **quickstart** field, type `pageviews`.
- 

1. Click **Continue**.

2. Review the connector configuration and click **Launch**.

MANAGEMENT ›

## Kafka Connect

Bring data in

Send data out

Search connectors

+ Add connector

Connectors			Details
Status	Name		Active tasks
Running	datagen-pag...	...	1

Run another instance of the [Kafka Connect Datagen](#) connector to produce Kafka data to the `users` topic in AVRO format.

Click **Add connector**.

Find the DatagenConnector tile and click **Connect**.

Name the connector `datagen-users`. After naming the connector, new fields appear. Scroll down and specify the following configuration values:

- Max Task : 1
- In the **Key converter class** field, type `org.apache.kafka.connect.storage.StringConverter`.
- In the **kafka.topic** field, type `users`.
- In the **max.interval** field, type `1000`.
- In the **iterations** field, type `1000000000`.
- In the **quickstart** field, type `users`.
  - Click **Continue**.
  - Review the connector configuration and click **Launch**.

At the end of this.

## Kafka Connect

[Bring data in](#)[Send data out](#)[+ Add connector](#)

### Connectors

Status

Name

### Details

Active tasks

Running

datagen-pag...

...

1

Running

datagen-users

...

1

Verify the messages in the both the topics:

Using the control centers:

Topics -> pageviews -> Messages:

The screenshot displays the Apache Kafka Control Center interface for the 'pageviews' topic. The left sidebar contains navigation links: Cluster overview, Brokers, Topics (highlighted), Connect, ksqlDB, Consumers, Replicators, and Cluster settings. Below these is a 'Health+' section with a 'New' button. The main content area is titled 'pageviews' and has tabs for Overview, Messages (selected), Schema, and Configuration. Under the 'Messages' tab, there are sections for Producers (Bytes in/sec: 2.68K) and Consumers (Bytes out/sec: 522). Below these is a 'Message fields' section listing topic, partition, offset, and timestamp. A 'Produce a new message to this topic' button is located above the message list. The message list shows two messages with details like Partition: 0, Offset: 7086, and Timestamp: 1641706995703. A 'Newest' button is visible on the right side of the message list.

## Topics -&gt; Users -&gt; Messages:

Cluster overview

Brokers

**Topics**

Connect

ksqlDB

Consumers

Replicators

Cluster settings

Health+ New

**Overview** **Messages** Schema Configuration

**Producers**

Bytes in/sec 295

**Consumers**

Bytes out/sec 5

**Message fields**

- topic
- partition
- offset
- timestamp
- timestampType
- headers
  - key

Filter by keyword

Jump to offset

offset

+ Produce a new message to this topic

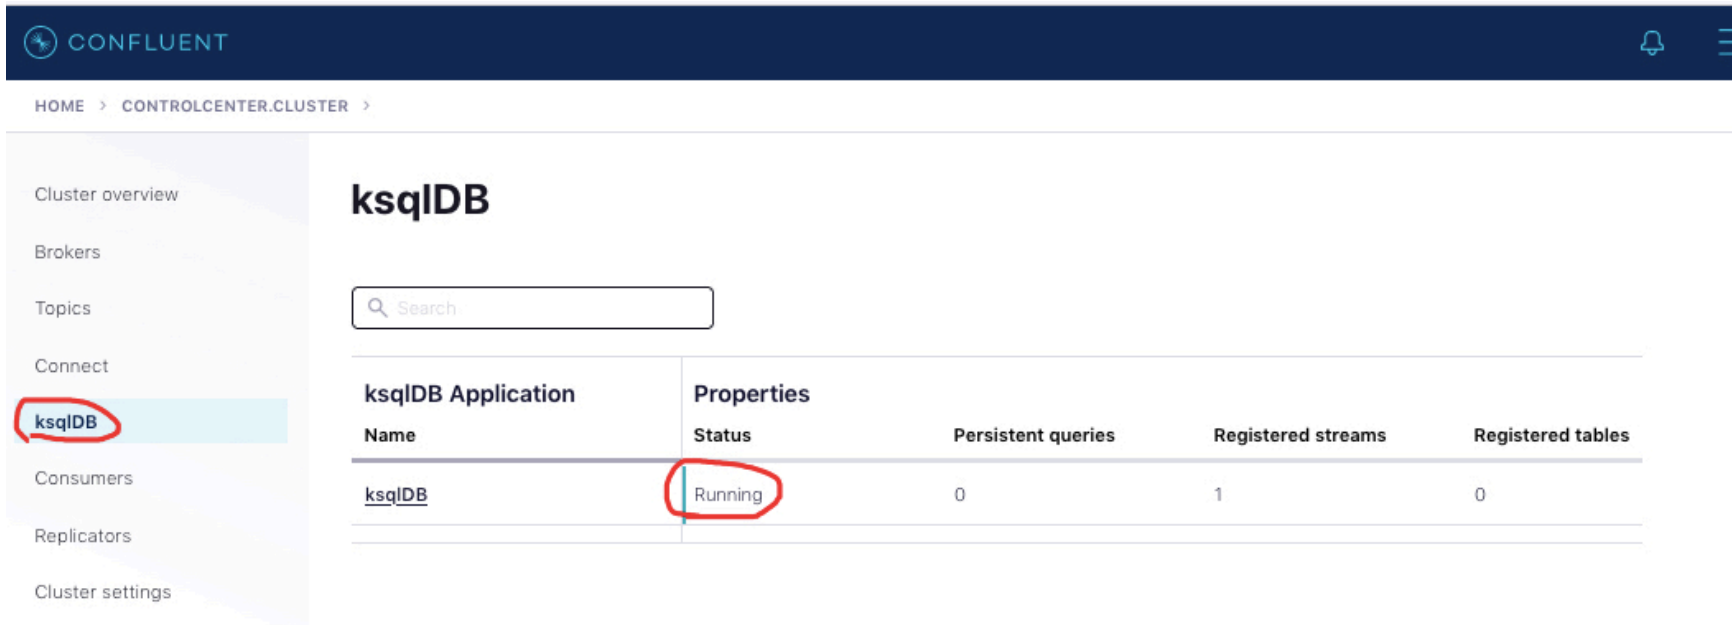
Newest

▼ {"registertime":1505601918721,"userid":"User\_8","regionid":"Region\_8","gender":"OTHER"}  
Partition: 0 Offset: 508 Timestamp: 1641707060914

▼ {"registertime":1512168635241,"userid":"User\_2","regionid":"Region\_6","gender":"FEMALE"}  
Partition: 0 Offset: 507 Timestamp: 1641707060763

▼ {"registertime":1487807651768,"userid":"User\_7","regionid":"Region\_2","gender":"MALE"}  
Partition: 0 Offset: 506 Timestamp: 1641707059920

Ensure that ksql DB services is up.



The screenshot shows the Confluent Control Center interface. On the left, a sidebar menu lists various cluster management options: Cluster overview, Brokers, Topics, Connect, **ksqlDB** (highlighted with a red circle), Consumers, Replicators, and Cluster settings. The main panel displays the 'ksqlDB' application details. At the top, there's a search bar. Below it, a table provides a summary of the application's status and configuration. The table has two main sections: 'ksqlDB Application' and 'Properties'. The 'ksqlDB Application' section includes a 'Name' column with the value 'ksqlDB'. The 'Properties' section includes columns for 'Status', 'Persistent queries', 'Registered streams', and 'Registered tables'. The 'Status' column shows 'Running', which is circled in red. The other columns show values: 0 for Persistent queries, 1 for Registered streams, and 0 for Registered tables.

ksqlDB Application		Properties		
Name	Status	Persistent queries	Registered streams	Registered tables
<a href="#">ksqlDB</a>	Running	0	1	0

If there is any issue, verify the status and configuration as shown below:

```
#confluent local services status
```



```
[root@ckafka0 ckafka]# confluent local services status
The local commands are intended for a single-node development environment only,
NOT for production usage. https://docs.confluent.io/current/cli/index.html

Using CONFLUENT_CURRENT: /opt/data/ckafka/confluent.652875
Connect is [UP]
Control Center is [UP]
Kafka is [UP]
Kafka REST is [UP]
ksqlDB Server is [UP]
Schema Registry is [UP]
ZooKeeper is [UP]
[root@ckafka0 ckafka]#
```

If unable to connect in 8088 port. Verify that the KSQL listeners IP and port are specify correctly in the configuration files.

/apps/confluent/etc/ksqldb/ksql-server.properties

listeners=http://localhost:8088 or

listeners=http://0.0.0.0:8088

Restart after any modification.

```
confluent local services ksql-server status
confluent local services ksql-server stop
confluent local services ksql-server start
confluent local services ksql-server status
After that verify the listening port.
```

lsof -i:8088

```
[root@ckafka0 ckafka]# lsof -i:8088
COMMAND  PID USER  FD   TYPE DEVICE SIZE/OFF NODE NAME
java     1092 root  628u  IPv4 140454      0t0  TCP localhost:37410->localhost:radan-http (ESTABLISHED)
java     1092 root  634u  IPv4 140457      0t0  TCP localhost:37414->localhost:radan-http (ESTABLISHED)
java     1092 root  637u  IPv4 145818      0t0  TCP localhost:37430->localhost:radan-http (ESTABLISHED)
java     1092 root  638u  IPv4 144459      0t0  TCP localhost:37432->localhost:radan-http (ESTABLISHED)
java     2968 root  502u  IPv4 143524      0t0  TCP localhost:radan-http (LISTEN)
java     2968 root  506u  IPv4 143555      0t0  TCP localhost:radan-http->localhost:37430 (ESTABLISHED)
java     2968 root  507u  IPv4 143556      0t0  TCP localhost:radan-http->localhost:37432 (ESTABLISHED)
java     2968 root  511u  IPv4 143551      0t0  TCP localhost:radan-http->localhost:37410 (ESTABLISHED)
java     2968 root  512u  IPv4 143552      0t0  TCP localhost:radan-http->localhost:37414 (ESTABLISHED)
```

It means, the KSQL server is running.

-----Lab Installation completes End here. -----

### 3. Workflow using KSQL - CLI – 90 Minutes

Following features will be demonstrated.

- Create Topics and Produce Data
- Create and produce data to the Kafka topics pageviews and users.
- Inspect Kafka Topics by Using SHOW and PRINT Statements
- Create a Stream and Table
- Write Queries

This tutorial demonstrates a simple workflow using KSQL to write streaming queries against messages in Kafka.

To get started, you must start a Kafka cluster, including ZooKeeper and a Kafka broker. KSQL will then query messages from this Kafka cluster. KSQL is installed in the Confluent Platform by default.

#### Prerequisites:

- [Confluent Platform](#) is installed and running. This installation includes a Kafka broker, KSQL, Control Center, ZooKeeper, Schema Registry, REST Proxy, and Connect.
- If you installed Confluent Platform via TAR or ZIP, navigate into the installation directory. The paths and commands used throughout this tutorial assume that you are in this installation directory.
- Consider [installing](#) the Confluent CLI to start a local installation of Confluent Platform.

- Java: Minimum version 1.8. Install Oracle Java JRE or JDK  $\geq 1.8$  on your local machine

### Create Topics and Produce Data

[https://ksqldb.io/quickstart.html?\\_ga=2.53841192.1438767497.1642131382-2002989446.1641377120&\\_gac=1.255954681.1642171371.CjwKCAiA24SPBhBoEiwAjBgkhg1qFCOJ-Ohq2cWlGrT9c323dWfPKKpOG6zXpZrNXjqUelgasqp5BoCTEoQAvD\\_BwE](https://ksqldb.io/quickstart.html?_ga=2.53841192.1438767497.1642131382-2002989446.1641377120&_gac=1.255954681.1642171371.CjwKCAiA24SPBhBoEiwAjBgkhg1qFCOJ-Ohq2cWlGrT9c323dWfPKKpOG6zXpZrNXjqUelgasqp5BoCTEoQAvD_BwE)

Create and produce data to the Kafka topics `pageviews` and `users`. These steps use the KSQL datagen that is included Confluent Platform.

1. Create the `pageviews` topic and produce data using the data generator. The following example continuously generates data with a value in DELIMITED format.

```
ksql-datagen quickstart=pageviews format=json topic=pageviews maxInterval=500
```

```
(base) [root@tos ~]#
(base) [root@tos ~]#
(base) [root@tos ~]# ksql-datagen quickstart=pageviews format=delimited topic=pa
geviews maxInterval=500
[2019-07-31 21:35:34,823] INFO AvroDataConfig values:
    schemas.cache.config = 1
    enhanced.avro.schema.support = false
    connect.meta.data = true
(io.confluent.connect.avro.AvroDataConfig:179)
1 --> ([ 1564589135082 | 'User_3' | 'Page_97' ]) ts:1564589135333
11 --> ([ 1564589135590 | 'User_7' | 'Page_66' ]) ts:1564589135591
21 --> ([ 1564589135857 | 'User_1' | 'Page_34' ]) ts:1564589135861
31 --> ([ 1564589135959 | 'User_6' | 'Page_37' ]) ts:1564589135959
41 --> ([ 1564589136036 | 'User_6' | 'Page_66' ]) ts:1564589136036
51 --> ([ 1564589136428 | 'User_2' | 'Page_98' ]) ts:1564589136428
61 --> ([ 1564589136761 | 'User_9' | 'Page_26' ]) ts:1564589136761
```

2. Produce Kafka data to the `users` topic using the data generator. The following example continuously generates data with a value in JSON format.

```
$ ksql-datagen quickstart=users format=json topic=users maxInterval=100
```

### Tip

You can also produce Kafka data using the `kafka-console-producer` CLI provided with Confluent Platform.

```
(base) [root@tos ~]#
(base) [root@tos ~]# ksql-datagen quickstart=pageviews format=delimited topic=pa
geviews maxInterval=500
[2019-07-31 21:35:34,823] INFO AvroDataConfig values:
    schemas.cache.config = 1
    enhanced.avro.schema.support = false
    connect.meta.data = true
(io.confluent.connect.avro.AvroDataConfig:179)
1 --> ([ 1564589135082 | 'User_3' | 'Page_97' ]) ts:1564589135333
11 --> ([ 1564589135590 | 'User_7' | 'Page_66' ]) ts:1564589135591
21 --> ([ 1564589135857 | 'User_1' | 'Page_34' ]) ts:1564589135861
31 --> ([ 1564589135959 | 'User_6' | 'Page_37' ]) ts:1564589135959
41 --> ([ 1564589136036 | 'User_6' | 'Page_66' ]) ts:1564589136036
51 --> ([ 1564589136428 | 'User_2' | 'Page_98' ]) ts:1564589136428
61 --> ([ 1564589136761 | 'User_9' | 'Page_26' ]) ts:1564589136761
```

## Launch the KSQL CLI

To launch the CLI, run the following command. It will route the CLI logs to the `./ksql_logs` directory, relative to your current directory. By default, the CLI will look for a KSQL Server running at `http://localhost:8088`.

```
$ LOG_DIR=./ksql_logs ksql
```

## Important

By default KSQL attempts to store its logs in a directory called `logs` that is relative to the location of the `ksql` executable. For example, if `ksql` is installed at `/usr/local/bin/ksql`, then it would attempt to store its logs in `/usr/local/logs`. If you are running `ksql` from the

default Confluent Platform location, `<path-to-confluent>/bin`, you must override this default behavior by using the `LOG_DIR` variable.

After KSQL is started, your terminal should resemble this.

[illegible]

## Inspect Kafka Topics By Using SHOW and PRINT Statements

KSQL enables inspecting Kafka topics and messages in real time.

- Use the `SHOW TOPICS` statement to list the available topics in the Kafka cluster.
- Use the `PRINT` statement to see a topic's messages as they arrive.

In the KSQL CLI, run the following statement:

```
SHOW TOPICS;
```

Your output should resemble:

Kafka Topic	Registered	Partitions	Partition Replicas	Consumers	ConsumerGroups
-----					
_confluent-metrics	false	12	1	0	0
_schemas	false	1	1	0	0
pageviews	false	1	1	0	0
users	false	1	1	0	0
-----					

Inspect the `users` topic by using the PRINT statement:

```
PRINT 'users';
```

Your output should resemble:

Format:JSON

```
{ "ROWTIME":1540254230041, "ROWKEY":"User_1", "registertime":1516754966866, "userid":"User_1", "regionid":"Region_9", "gender":"MALE" }
```



```
{ "ROWTIME":1540254230081,"ROWKEY":"User_3","registertime":1491558386780,"userid":"User_3","regionid":"Region_2","gender":"MALE"}
{"ROWTIME":1540254230091,"ROWKEY":"User_7","registertime":1514374073235,"userid":"User_7","regionid":"Region_2","gender":"OTHER"}
^C{"ROWTIME":1540254232442,"ROWKEY":"User_4","registertime":1510034151376,"userid":"User_4","regionid":"Region_8","gender":"FEMALE"}
Topic printing ceased
```

Press CTRL+C to stop printing messages.

Inspect the `pageviews` topic by using the PRINT statement:

```
PRINT 'pageviews';
```

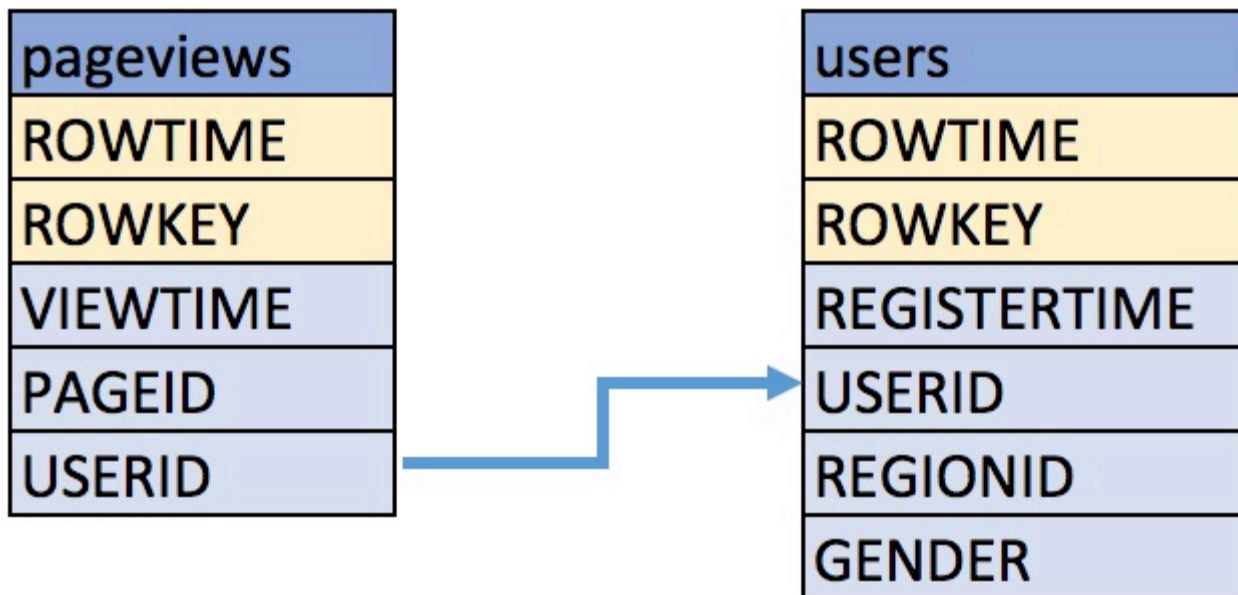
Your output should resemble:

```
Format:STRING
10/23/18 12:24:03 AM UTC , 9461 , 1540254243183,User_9,Page_20
10/23/18 12:24:03 AM UTC , 9471 , 1540254243617,User_7,Page_47
10/23/18 12:24:03 AM UTC , 9481 , 1540254243888,User_4,Page_27
^C10/23/18 12:24:05 AM UTC , 9521 , 1540254245161,User_9,Page_62
Topic printing ceased
ksql>
```

Press CTRL+C to stop printing messages.

Create a Stream and Table

These examples query messages from Kafka topics called `pageviews` and `users` using the following schemas:



1. Create a stream, named `pageviews_original`, from the `pageviews` Kafka topic, specifying the `value_format` of `DELIMITED`.

```
CREATE STREAM pageviews_original (viewtime bigint, userid varchar, pageid varchar) WITH  
ITH  
(kafka_topic='pageviews', value_format='JSON');
```

Your output should resemble:

```
ksql> CREATE STREAM pageviews_original (viewtime bigint, userid varchar, pageid varchar) WITH  
>(kafka_topic='pageviews', value_format='DELIMITED');  
>  
  
Message  
-----  
Stream created  
-----  
ksql> █
```

### Tip

You can run `DESCRIBE pageviews_original;` to see the schema for the stream. Notice that KSQL created two additional columns, named `ROWTIME`, which corresponds with the Kafka message timestamp, and `ROWKEY`, which corresponds with the Kafka message key.

```
ksql> DESCRIBE pageviews_original;

Name          : PAGEVIEWS_ORIGINAL
Field  | Type
-----|-----
ROWTIME | BIGINT      (system)
ROWKEY  | VARCHAR(STRING) (system)
VIEWTIME | BIGINT
USERID  | VARCHAR(STRING)
PAGEID  | VARCHAR(STRING)
-----|-----

For runtime statistics and query details run: DESCRIBE EXTENDED <Stream,Table>;
ksql>
```

2. Create a table, named `users_original`, from the `users` Kafka topic, specifying the `value_format` of `JSON`.

```
CREATE TABLE users_original (registertime BIGINT, gender VARCHAR, regionid VARCHAR,
userid VARCHAR PRIMARY KEY) WITH
(kafka_topic='users', value_format='JSON');
```

Your output should resemble:

Message

-----  
Table created  
-----

**Tip**

You can run `DESCRIBE users_original;` to see the schema for the Table.

3. Optional: Show all streams and tables.

```
ksql> SHOW STREAMS;
```

Stream Name	Kafka Topic	Format
PAGEVIEWS_ORIGINAL	pageviews	DELIMITED

```
ksql> SHOW TABLES;
```

Table Name	Kafka Topic	Format	Windowed
USERS_ORIGINAL	users	JSON	false

## Write Queries

```
SET 'auto.offset.reset'='earliest';
```

These examples write queries using KSQL.

**Note:** By default KSQL reads the topics for streams and tables from the latest offset.

1. Use **SELECT** to create a query that returns data from a STREAM. This query includes the **LIMIT** keyword to limit the number of rows returned in the query result. Note that exact data output may vary because of the randomness of the data generation.

```
SELECT pageid FROM pageviews_original EMIT changes LIMIT 3;
```

Your output should resemble:

```
Page_24  
Page_73  
Page_78  
LIMIT reached  
Query terminated
```

2. Create a persistent query by using the **CREATE STREAM** keywords to precede the **SELECT** statement. The results from this query are written to the **PAGEVIEWS\_ENRICHED** Kafka topic. The following query enriches the **pageviews\_original** STREAM by doing a **LEFT JOIN** with the **users\_original** TABLE on the user ID.

```
CREATE STREAM pageviews_enriched AS  
SELECT users_original.userid AS userid, pageid, regionid, gender  
FROM pageviews_original
```

**LEFT JOIN users\_original**

**ON pageviews\_original.userid = users\_original.userid**

**Emit changes;**

Your output should resemble:

Message

-----

Stream created **and** running

-----

### Tip

You can run **DESCRIBE pageviews\_enriched;** to describe the stream.

3. Use **SELECT** to view query results as they come in. To stop viewing the query results, press **<ctrl-c>**. This stops printing to the console but it does not terminate the actual query. The query continues to run in the underlying KSQL application.

4. **SELECT \* FROM pageviews\_enriched Emit Changes;**

Your output should resemble:

IUser_9	IPage_92	IRegion_2	IMALE	
IUser_2	IPage_66	IRegion_6	IMALE	
IUser_3	IPage_10	IRegion_7	IMALE	
IUser_5	IPage_30	IRegion_3	IOTHER	
IUser_2	IPage_85	IRegion_6	IMALE	
IUser_1	IPage_46	IRegion_7	IOTHER	
IUser_6	IPage_56	IRegion_3	IFEMALE	
IUser_8	IPage_13	IRegion_2	IMALE	
IUser_4	IPage_19	IRegion_4	IFEMALE	
IUser_3	IPage_44	IRegion_7	IMALE	
IUser_8	IPage_57	IRegion_2	IMALE	
IUser_8	IPage_39	IRegion_2	IMALE	
IUser_9	IPage_15	IRegion_2	IMALE	
IUser_9	IPage_71	IRegion_2	IMALE	
IUser_7	IPage_69	IRegion_8	IMALE	

5. Create a new persistent query where a condition limits the streams content, using **WHERE**. Results from this query are written to a Kafka topic called **PAGEVIEWS\_FEMALE**.

```
CREATE STREAM pageviews_female AS
SELECT * FROM pageviews_enriched
WHERE gender = 'FEMALE';
```

Your output should resemble:

Message

-----  
Stream created **and** running  
-----



**Tip**

You can run `DESCRIBE pageviews_female;` to describe the stream.

6. Create a new persistent query where another condition is met, using `LIKE`. Results from this query are written to the `pageviews_enriched_r8_r9` Kafka topic.

```
CREATE STREAM pageviews_female_like_89
  WITH (kafka_topic='pageviews_enriched_r8_r9') AS
  SELECT * FROM pageviews_female
  WHERE regionid LIKE '%_8' OR regionid LIKE '%_9';
```

Your output should resemble:

Message

-----  
Stream created **and** running  
-----

7. Verify the above 2 streams:
 

```
select * from PAGEVIEWS_FEMALE_LIKE_89 emit changes limit 6;
select * from PAGEVIEWS_FEMALE emit changes limit 3;
```

```
ksql> select * from PAGEVIEWS_FEMALE_LIKE_89 emit changes limit 6;
```

USERID	PAGEID	REGIONID	GENDER
User_9	Page_15	Region_9	FEMALE
User_9	Page_17	Region_8	FEMALE
User_9	Page_66	Region_8	FEMALE
User_9	Page_62	Region_8	FEMALE
User_9	Page_71	Region_8	FEMALE
User_6	Page_31	Region_8	FEMALE

```
Limit Reached
Query terminated
ksql> select * from PAGEVIEWS_FEMALE emit changes limit 3;
```

USERID	PAGEID	REGIONID	GENDER
User_1	Page_30	Region_8	FEMALE
User_3	Page_23	Region_6	FEMALE
User_1	Page_81	Region_8	FEMALE

```
Limit Reached
Query terminated
ksql>
```

8. Create a new persistent query that counts the pageviews for each region combination in a **tumbling window** of 30 seconds when the count is greater than one. Results from this query are written to the **PAGEVIEWS\_REGIONS** Kafka topic in the Avro format. **KSQL**

will register the Avro schema with the configured Schema Registry when it writes the first message to the `PAGEVIEWS_REGIONS` topic.

```
CREATE TABLE pageviews_regions
WITH (
  KAFKA_TOPIC = 'pageviews_regions',VALUE_FORMAT='AVRO'
) AS
SELECT regionid , COUNT(*) AS numusers
FROM pageviews_enriched
WINDOW TUMBLING (size 30 second)
GROUP BY regionid
HAVING COUNT(*) > 1 emit changes;
```

Your output should resemble:

Message

-----  
Table created **and** running  
-----

### Tip

You can run `DESCRIBE pageviews_regions;` to describe the table.

9. Optional: View results from the above queries using `SELECT`.

```
SELECT regionid, numusers FROM pageviews_regions LIMIT 5;
```

Your output should resemble:

```
ksql> SELECT regionid, numusers FROM pageviews_regions emit changes LIMIT 5;
+-----+-----+
|REGIONID|NUMUSERS|
+-----+-----+
|Region_2|221     |
|Region_3|6169    |
|Region_5|10659   |
|Region_2|11476   |
|Region_9|2259    |
Limit Reached
Query terminated
```

10. Optional: Show all persistent queries.

11. **SHOW QUERIES;**

Your output should resemble:

Query ID	Kafka Topic	Query String
-----	-----	-----
-----	-----	-----

```

-----
-----
CSAS_PAGEVIEWS_FEMALE_1 | PAGEVIEWS_FEMALE | CREATE STREA
M pageviews_female AS SELECT * FROM pageviews_enriched WHERE gender =
'FEMALE';
CTAS_PAGEVIEWS_REGIONS_3 | PAGEVIEWS_REGIONS | CREATE TABLE
pageviews_regions WITH (VALUE_FORMAT='avro') AS SELECT gender, region
id , COUNT(*) AS numusers FROM pageviews_enriched WINDOW TUMBLING
(size 30 second) GROUP BY gender, regionid HAVING COUNT(*) > 1;
CSAS_PAGEVIEWS_FEMALE_LIKE_89_2 | PAGEVIEWS_FEMALE_LIKE_89 | CRE
ATE STREAM pageviews_female_like_89 WITH (kafka_topic='pageviews_enriche
d_r8_r9') AS SELECT * FROM pageviews_female WHERE regionid LIKE '%_8' O
R regionid LIKE '%_9';
CSAS_PAGEVIEWS_ENRICHED_o | PAGEVIEWS_ENRICHED | CREATE STR
EAM pageviews_enriched AS SELECT users_original.userid AS userid, pageid, regio
nid, gender FROM pageviews_original LEFT JOIN users_original ON pagevie
ws_original.userid = users_original.userid;
-----
-----
-----
-----

```

For detailed information on a Query run: EXPLAIN <Query ID>;

12. Optional: Examine query run-time metrics and details. Observe that information including the target Kafka topic is available, as well as throughput figures for the messages being processed.

```
DESCRIBE PAGEVIEWS_REGIONS EXTENDED;
```

Your output should resemble:

```
Name       : PAGEVIEWS_REGIONS
Type       : TABLE
Key field   : KSQL_INTERNAL_COL_o|+|KSQL_INTERNAL_COL_1
Key format  : STRING
Timestamp field : Not set - using <ROWTIME>
Value format : AVRO
Kafka topic  : PAGEVIEWS_REGIONS (partitions: 4, replication: 1)
```

Field | Type

```
-----
ROWTIME | BIGINT      (system)
ROWKEY   | VARCHAR(STRING) (system)
GENDER   | VARCHAR(STRING)
REGIONID | VARCHAR(STRING)
NUMUSERS | BIGINT
-----
```

Queries that write into this TABLE

```
-----
CTAS_PAGEVIEWS_REGIONS_3 : CREATE TABLE pageviews_regions WITH (value_format='avro') AS
SELECT gender, regionid , COUNT(*) AS numusers FROM
```

```
pageviews_enriched    WINDOW TUMBLING (size 30 second)    GROUP BY gender,
regionid    HAVING COUNT(*) > 1;
```

For query topology **and** execution plan please run: EXPLAIN <QueryId>

Local runtime statistics

```
-----
messages-per-sec:  3.06  total-messages:  1827  last-message: 7/19/18 4:17:55 PM
UTC
failed-messages:   0  failed-messages-per-sec:   0  last-failed:  n/a
(Statistics of the local KSQL server interaction with the Kafka topic PAGEVIEWS_REGI
ONS)
ksql>
```

----- Lab Ends Here -----