

Stream v.s. Table?


```
KStream<..> stream1 = builder.stream("topic1");
```

```
KStream<..> stream2 = builder.stream("topic2");
```

```
KStream<..> joined = stream1.leftJoin(stream2, ...);
```

```
KTable<..> aggregated = joined.aggregateByKey(...);
```

```
aggregated.to("topic2");
```



Tables \approx *Streams*

TABLES \approx STREAMS

(key1, value1)

(key2, value2)

(key1, value3)

•

•

•

TABLES \approx STREAMS

(key1, value1) \rightarrow

key1	value1
------	--------

(key2, value2) \rightarrow

key1	value1
key2	value2

(key1, value3) \rightarrow

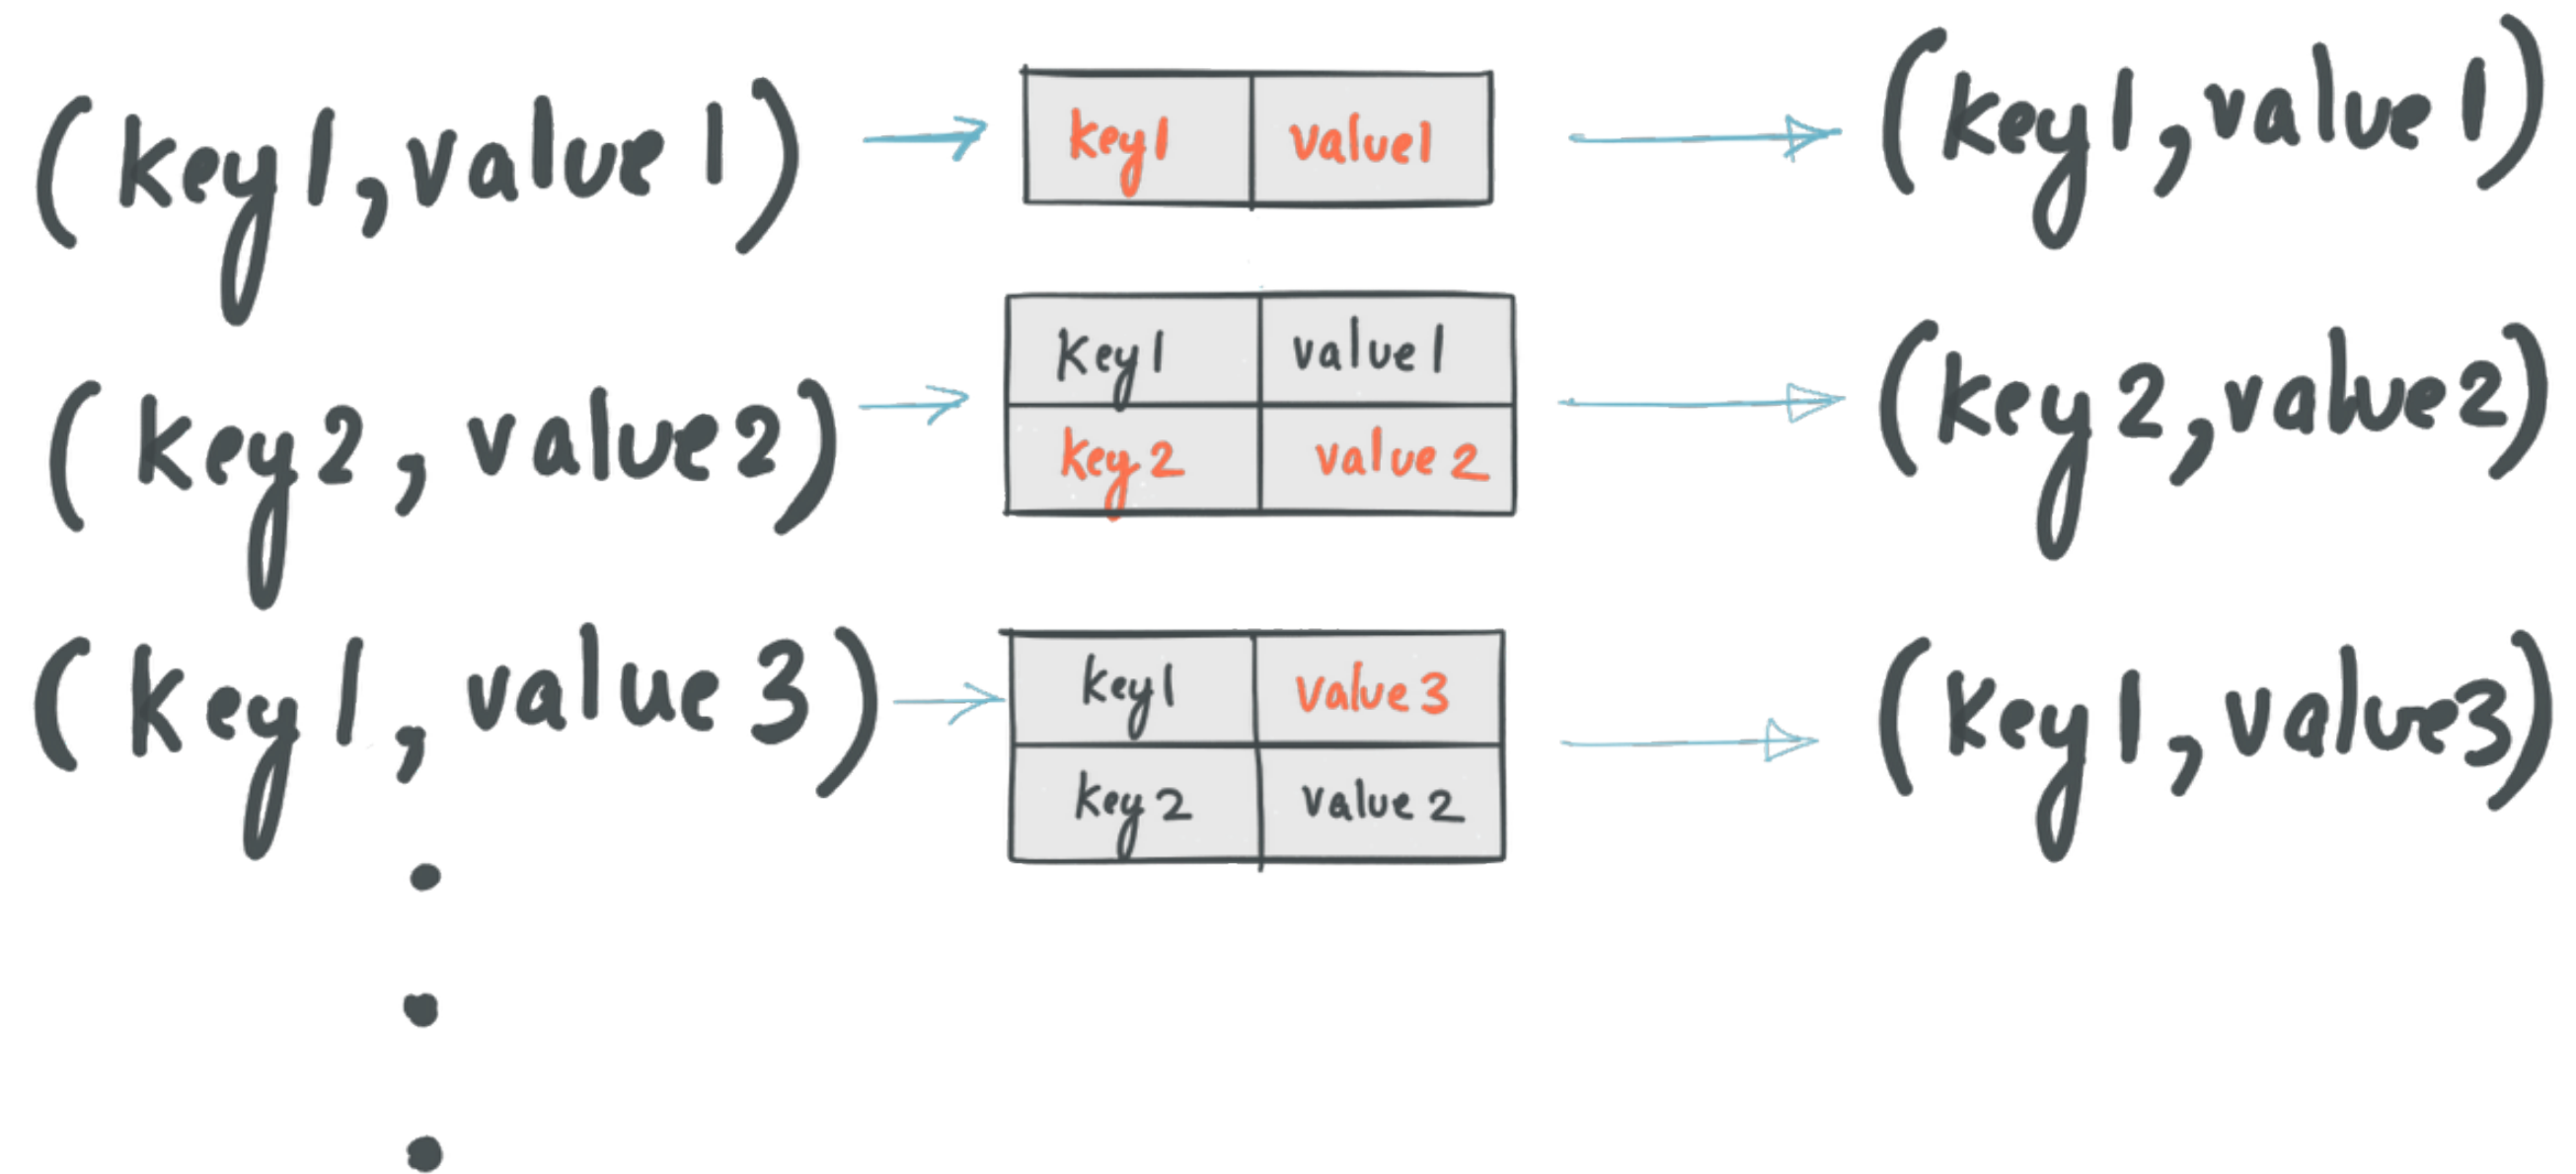
key1	value3
key2	value2

•

•

•

TABLES \approx STREAMS



The Stream-Table Duality

- *A **stream** is a changelog of a **table***
- *A **table** is a materialized view at time of a **stream***
- *Example: change data capture (CDC) of databases*

KStream = *interprets data as record stream*

~ think: “append-only”

KTable = *data as changelog stream*

~ continuously updated materialized view

KStream

User purchase history



KTable

User employment profile



KStream

User purchase history



"Alice bought **eggs**."

KTable

User employment profile



"Alice is now at **LinkedIn**."

KStream

User purchase history



"Alice bought **eggs** and **milk**."

KTable

User employment profile



"Alice is now at LinkedIn
Microsoft."

KStream.aggregate()

time

(key: Alice, value: 2)



KTable.aggregate()

(key: Alice, value: 2)

KStream.aggregate()

time

(key: Alice, value: **2+3**)



KTable.aggregate()

(key: Alice, value: ~~2~~ **3**)

reduce()
aggregate()
...

...

map()
filter()
join()
...


KStream

KTable


map()
filter()
join()
...

toStream()


Updates Propagation in KTable




Updates Propagation in KTable



Updates Propagation in KTable



Updates Propagation in KTable



***Lab: Running your first Kafka Streams Application:
WordCount – 60 minutes***